

Performance Monitoring AlwaysOn Availability Groups

Anthony E. Nocentino
aen@centinosystems.com



Anthony E. Nocentino

- **Consultant and Trainer**
- **Founder and President of Centino Systems**
 - Specialize in system architecture and performance
 - Masters Computer Science (almost a PhD)
 - Friend of Redgate - 2015/2016
 - Microsoft Certified Professional
- **email:** aen@centinosystems.com
- **Twitter:** @nocentino
- **Blog:** www.centinosystems.com/blog
- **Pluralsight Author:** www.pluralsight.com



Overview

- Motivation
- How availability groups move data
- Impact of replication latency on availability
- Monitoring techniques
- Demo
- Dealing with replication latency

Why is this important?

- Recovery Objectives
 - Recovery Point Objective - RPO
 - Recovery Time Objective - RTO
- Availability
 - How much data can we lose?
 - How fast will the system fail over?
- Monitoring and Trending
 - Establish a baseline for analysis - are we meeting those objectives?
 - Impact on resources
- Ownership
 - All of the components are monitored by the DBA

Data Movement In Availability Groups

- Transaction log blocks are replicated to secondaries
- Replication mode
 - Synchronous
 - Asynchronous
- Database mirroring endpoint

Network Based Replication

- Strong working relationship with network team
 - Maintenance - patching, network outages, database
- Network conditions can impact your AG's availability
 - Latency - how **long** it takes for a packet of data to traverse the network from source to destination.
 - Bandwidth - how **much** data can be moved in a time interval

Network Latency

- Often measured in milliseconds, sometimes microseconds
- Directly impacts network throughput
- TCP sliding window
- ping isn't your best measure of latency, by default it doesn't include any load...measure **your** workload
- It's often up to us to **PROVE** to the network team there is an issue
 - **Pinging 192.168.2.1 with 32 bytes of data:**
 - **Reply from 192.168.2.1: bytes=32 time=1001ms TTL=128**
- In synchronous mode, you have to wait for network latency

Availability Group Flow Control

- Used in response to network and system conditions
- Log blocks exchange sequence numbers
- The AG will enter flow control mode IF:
 - The **primary** detects too many unacknowledged messages, the primary stops sending messages
 - The **secondary** needs to tell the primary to back off, likely due to resource constraints, it will send a flow control message to the primary to back off
- Primary polls every 1000ms for a change in flow control state
- Secondary will message primary to leave flow control mode

From: SQL Server PFE Blog - <http://bit.ly/1ZpGyIL>

Database Synchronization States

- Not synchronizing
- Synchronized
- Synchronizing
- Reverting
- Initializing

<https://msdn.microsoft.com/en-us/library/ff877972.aspx>

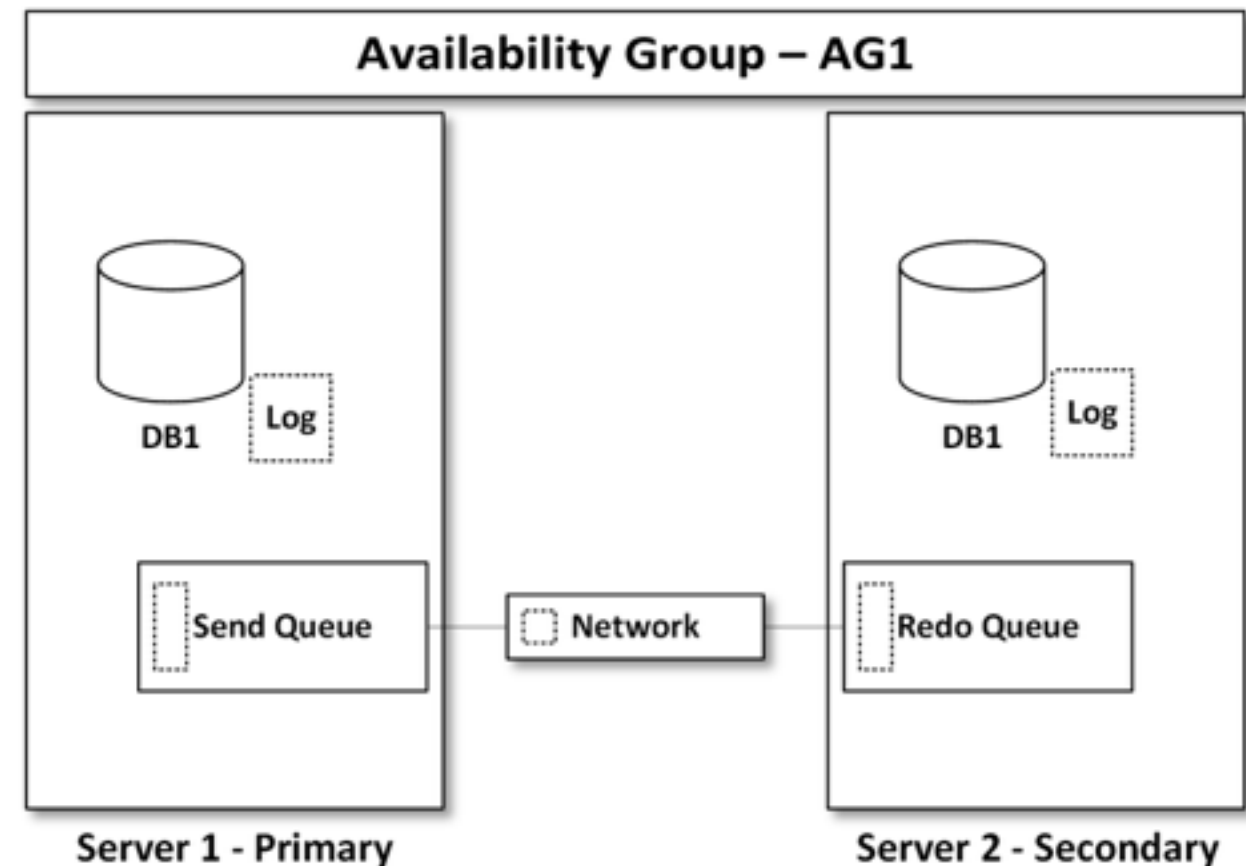
Failover Modes

- Automatic
 - Synchronous mode only
 - Commonly used within a data center
 - Synchronization state must be synchronized
- Manual
 - Synchronous or Asynchronous
 - Commonly used between data centers

<https://msdn.microsoft.com/en-us/library/hh213151.aspx>

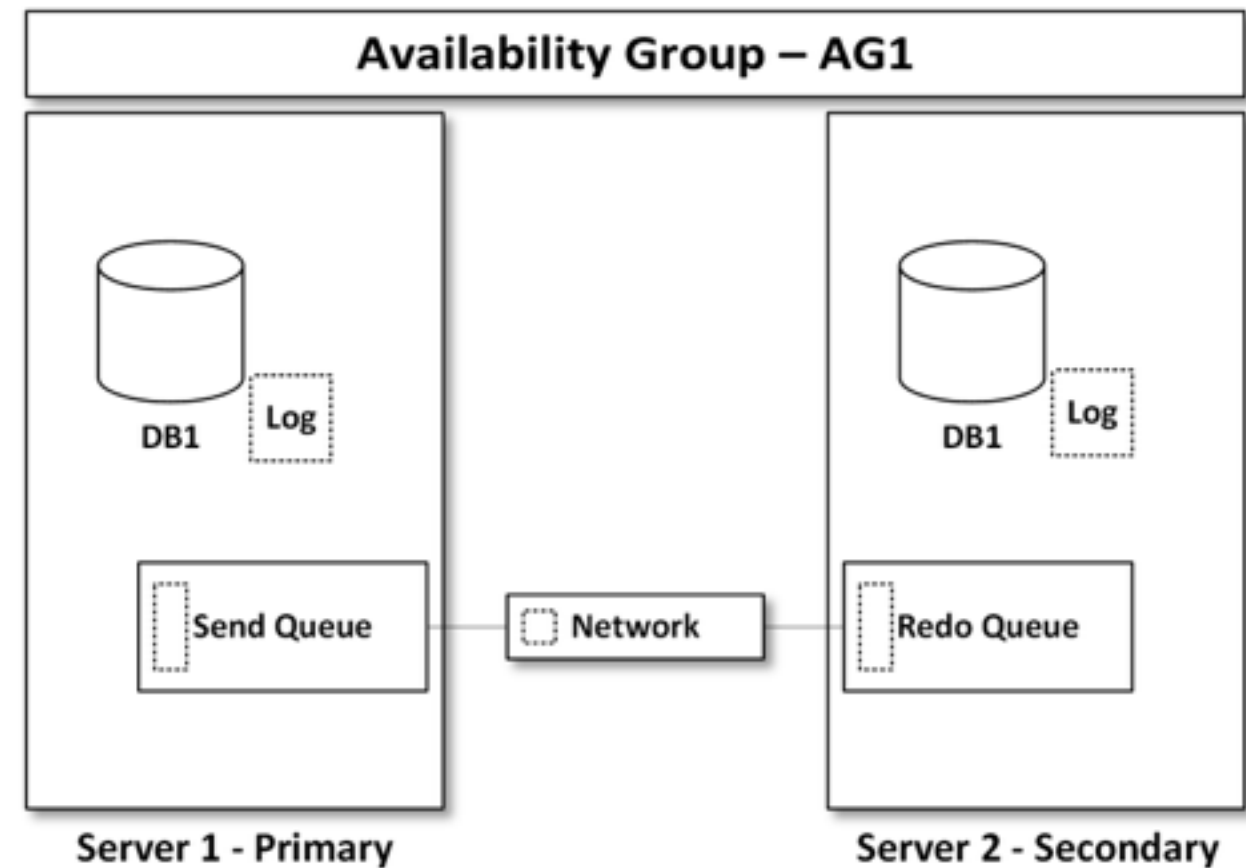
Send Queue

- Queues log blocks to be sent to the secondaries
- Each replica maintains it's own view of the send queue
- Queued data is at risk to data loss in the event of a primary failure
- The send queue can grow due to an unreachable secondary, network outage, network latency and large amount of data change



Redo Queue

- Queues log blocks received on the secondary
- Each replica has its own redo queue
- On failover, the redo queue must be completely processed
- The redo queue can grow due to a slow disk subsystem or resource contention or sustained outage and subsequent reconnection of a secondary




Send Queue Impact on Availability

- When log generation on primary exceeds the rate they can be sent to the secondaries...
 - No automatic failover
 - Data loss
 - Stale data for reporting from secondaries
 - Stale data for off-loaded backups on secondaries
 - Off-loaded log backups can fail
 - Transaction delay
 - Fill up transaction logs
- **Even in synchronous mode!**

Redo Queue Impact on Availability

- When log blocks received on the secondary exceed the rate they can be processed by the redo thread...
 - Delayed failover
 - Detect failure
 - Process Redo Queue
 - Crash recover database
 - Stale data for reporting from secondaries
 - Stale data for off-loaded backups on secondaries
 - Off-loaded log backups can fail
 - Transaction delay

Log Backups



Transaction Delay

- In synchronous mode, when secondaries are behind, queries on the primary can be delayed
- HADR_SYNC_COMMIT
- HADR_SYNCHRONIZING_THROTTLE - replica back online

Maintenance Events That Can Impact Availability

- Bulk data modifications
 - Database maintenance
 - Network or server maintenance
-
- Carefully plan maintenance
 - Collaborate with other teams!

Monitoring AG Performance

- Dynamic Management Views
 - `sys.dm_hadr_database_replica_states`
- Perfmon Counters
 - SQL Server:Availability Replica
 - Replication data - messages sent, bytes sent, flow control
 - SQL Server:Database Replica
 - Database data - log bytes sent, queue sizes, transaction delay per database

Measuring Replication Latency

- `sys.dm_hadr_database_replica_states`
 - `log_send_queue_size`
 - `log_send_rate`
 - `redo_queue_size`
 - `redo_queue_rate`
- On the primary there's a row for each database on **each** replica
- On the secondaries there's a row for each database on **that** replica
- Replicas track their own values
- When a replica goes offline...
 - `log_send_queue_size` changes to NULL

log send queue is from primary to secondary

Measuring Replication Latency - **ugh!!!**

- **Well, it looks like `sys.dm_hadr_database_replica_states` doesn't report the correct values for `log_send_rate` and `redo_queue_rate`**
 - Documented as KB
 - Reported on Connect
 - <https://connect.microsoft.com/SQLServer/Feedback/Details/928582>
 - Known bug in SQL Server 2012 or 2014
 - <https://support.microsoft.com/en-us/kb/3012182>
 - Cumulative Update 5 or better
 - Observed in SQL 2016
 - Perfmon!

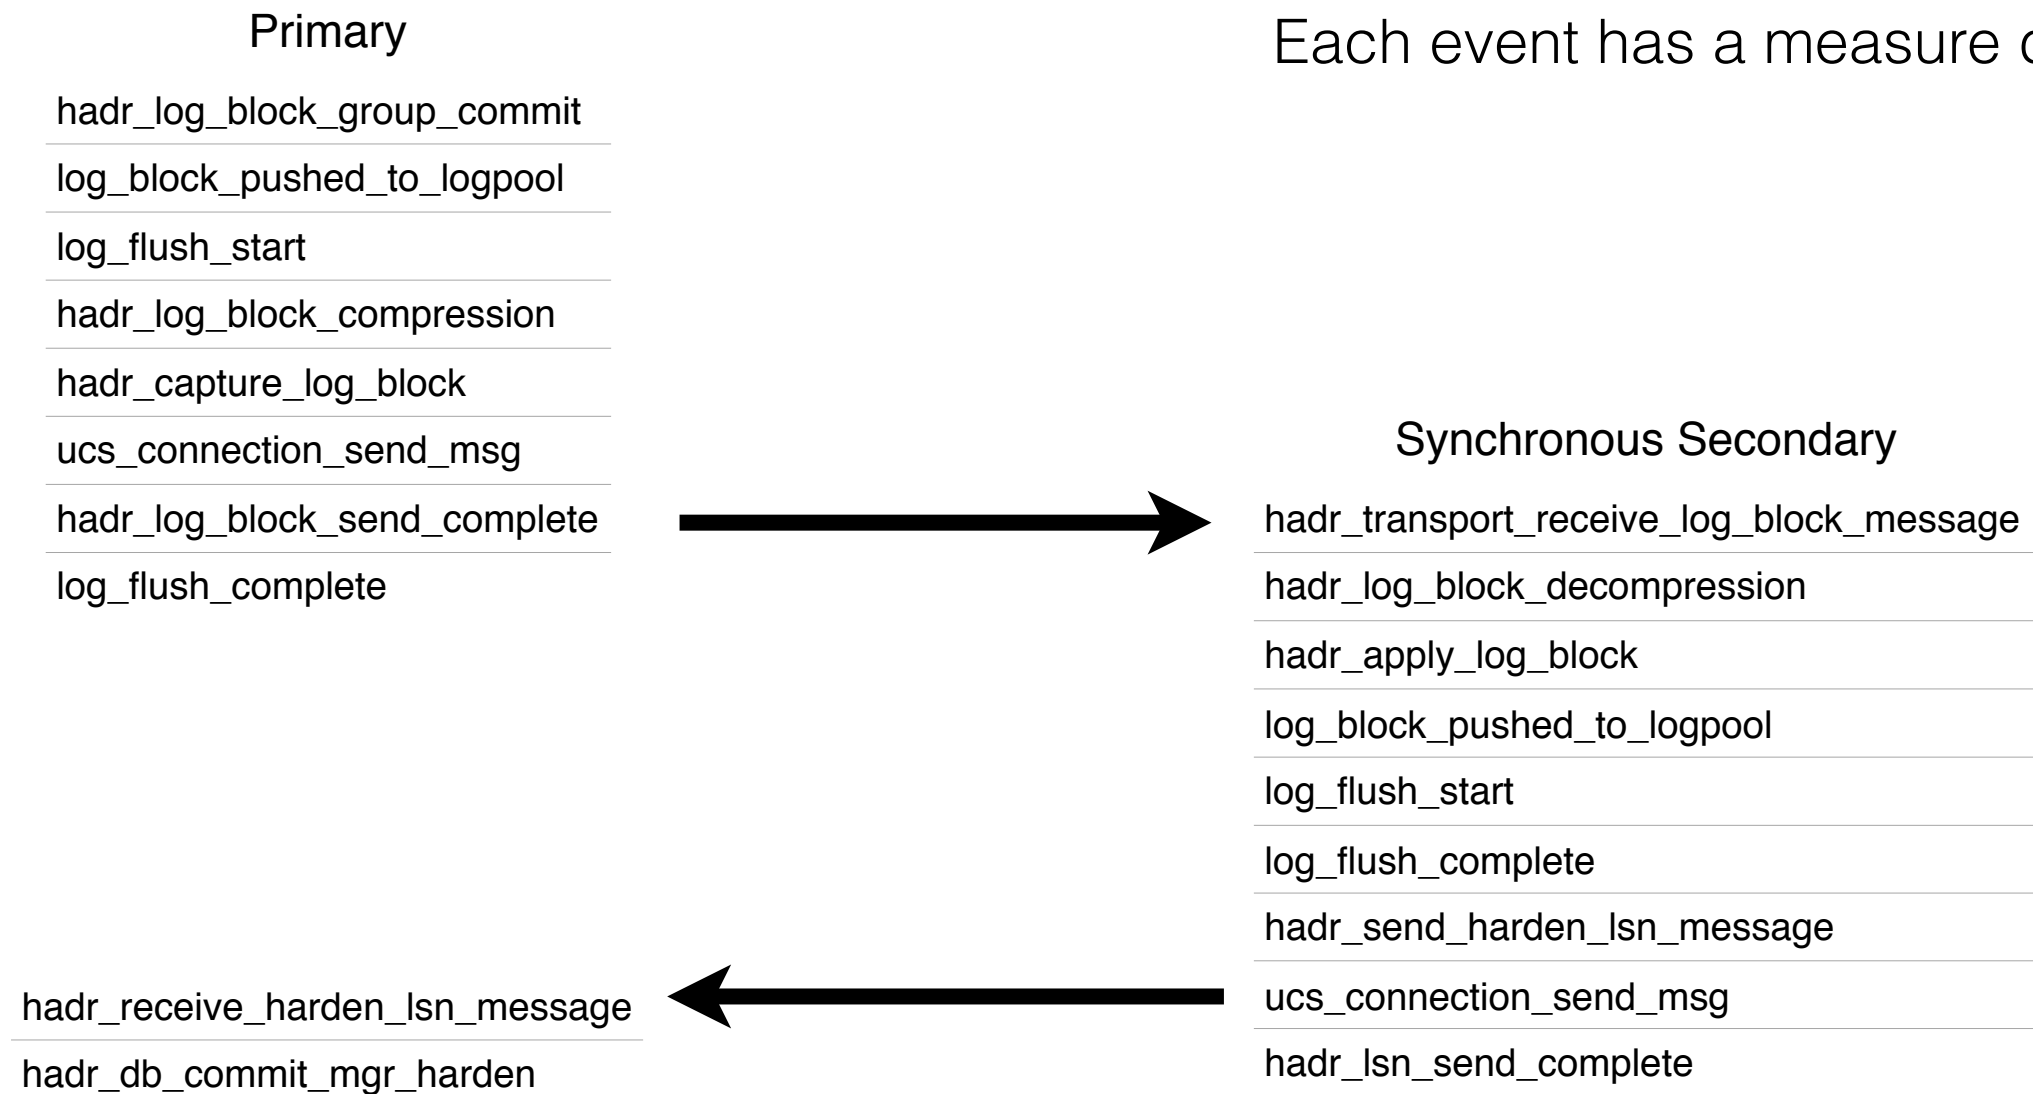
Measuring Latency with Perfmon

- Primary
 - SQLServer:Databases - Log Bytes Flushed/sec
 - SQLServer:Availability Replica - Bytes Sent to Replica/sec (compressed)
 - Network Interface - Bytes Sent/sec
- Secondaries
 - SQLServer:Availability Replica - Bytes Received From Replica (compressed)
 - SQLServer:Database Replica - Log Bytes Received/sec (log send rate/decompressed)
 - SQLServer:Database Replica - Redone Bytes/sec (log redo rate)
 - Network Interface - Bytes Received/sec

Measuring Latency with Extended Events

In SQL 2014 - SP2

Each event has a measure duration



Wait stats - sync vs. async

Synchronous - HADR_SYNC_COMMIT

```
sqldk.dll!XeSosPkg::wait_info::Publish+0x138
sqldk.dll!SOS_Scheduler::UpdateWaitTimeStats+0x2bc
sqldk.dll!SOS_Task::PostWait+0x9e
sqlmin.dll!EventInternal<SuspendQueueSLock>::Wait+0x1fb
sqlmin.dll!SequencedObject<LogBlockId,SequencedWaitInfo<LogBlockId>0>::WaitUntilSequenceAdvances+0x160
sqlmin.dll!HaDrCommitMgr::HardenNotifyInternal+0x1af
sqlmin.dll!HaDrCommitMgr::HardenNotify+0xac
sqlmin.dll!RecoveryUnit::NotifyHardenParticipants+0x1e2
sqlmin.dll!RecoveryUnit::HardenLog+0x217
sqlmin.dll!XdesRMFull::CommitInternal+0x6b8
sqlmin.dll!XactRM::SinglePhaseCommit+0x1a1
sqlmin.dll!XactRM::CommitInternal+0x472
...
```


Asynchronous - WRITELOG

```
sqldk.dll!XeSosPkg::wait_info::Publish+0x138
sqldk.dll!SOS_Scheduler::UpdateWaitTimeStats+0x2bc
sqldk.dll!SOS_Task::PostWait+0x9e
sqlmin.dll!SQLServerLogMgr::WaitLCFlush+0x219
sqlmin.dll!SQLServerLogMgr::LogFlush+0x29e
sqlmin.dll!SQLServerLogMgr::WaitLogWritten+0x17
sqlmin.dll!RecoveryUnit::HardenLog+0x25e
sqlmin.dll!XdesRMFull::CommitInternal+0x6b8
sqlmin.dll!XactRM::SinglePhaseCommit+0x1a1
sqlmin.dll!XactRM::CommitInternal+0x472
...
```

Monitoring Tools


- Build your own
- AlwaysOn Dashboard
- Third Party Tool
 - SQL Sentry Performance Advisor
 - Redgate SQL Monitor

Demo




Demo

Demo



Real World Example



Dealing With Slow Replication Latency

- Identify your bottleneck and mitigate it
 - Minimize log generation
 - Use smart index maintenance/Better Indexes
 - More bandwidth
 - Perhaps a dedicated network connection
 - Better hardware
 - Log throughput on secondaries needs to be equal to primary
- Upgrade SQL Server
 - 2012 single threaded redo - ~45MB/sec
 - 2016 multi-threaded redo - ~600MB/sec

Key Takeaways

- It is imperative to track and trend replication latency in your Availability Groups so you can answer the questions
 - How much data can I lose?
 - How long it will take to failover?
- Monitor and trend `send_queue` and `redo_queue` in `sys.dm_hadr_database_replica_states` on replicas to measure availability impact
- Understand how much log is generated in your databases
- Understand your system's operations, consider downtime for patching and network maintenance

Key Takeaways

- Plan database maintenance
- Use a smart index maintenance strategy!
- Offloaded backups
 - If availability is most important, backup on primary

Need more data?

<http://www.centinosystems.com/blog/talks/>

Links to resources

Demos

Presentation

aen@centinosystems.com



Free SQL Monitor!

THIS IS AN AMAZING DEAL WORTH \$3000!!!

**Thanks for Attending
SQL Saturday Baton Rouge 2016!**

- Speaker evaluations: Use the small square cards at the front of the classroom, give directly to speaker
- Speaker: Please give out 1 book ticket
- Book Ticket Winner: Bring your ticket to the user group booth in the main atrium to redeem (supplies limited)

SQL Saturday Baton Rouge 2016 Sponsors

Platinum Level



Gold Level



Silver Level



Bronze Level



Blog Level



Our SQL Saturday Host



#515 | BATON ROUGE 2016



Questions?

References

- <http://www.centinosystems.com/blog/sql/designing-for-offloaded-backups-in-alwayson-availability-groups/>
- <http://www.centinosystems.com/blog/sql/designing-for-offloaded-log-backups-in-alwayson-availability-groups-monitoring/>
- <http://www.centinosystems.com/blog/sql/monitoring-availability-groups-with-redgates-sql-monitor>
- <https://msdn.microsoft.com/en-us/library/ff878537.aspx>
- <https://msdn.microsoft.com/en-us/library/ff877972.aspx>