

# Performance Monitoring AlwaysOn Availability Groups

Anthony E. Nocentino  
[aen@centinosystems.com](mailto:aen@centinosystems.com)



# Anthony E. Nocentino

- **Consultant and Trainer**
- **Founder and President of Centino Systems**
  - Specialize in system architecture and performance
  - Masters Computer Science (almost a PhD)
  - Microsoft MVP - Data Platform - 2017
  - Friend of Redgate - 2015/2016
- **email:** [aen@centinosystems.com](mailto:aen@centinosystems.com)
- **Twitter:** @nocentino
- **Blog:** [www.centinosystems.com/blog](http://www.centinosystems.com/blog)
- **Pluralsight Author:** [www.pluralsight.com](http://www.pluralsight.com)



# Overview

- Motivation
- How Availability Groups move data
- Impact of replication latency on availability
- Monitoring techniques
- Demo
- Dealing with replication latency

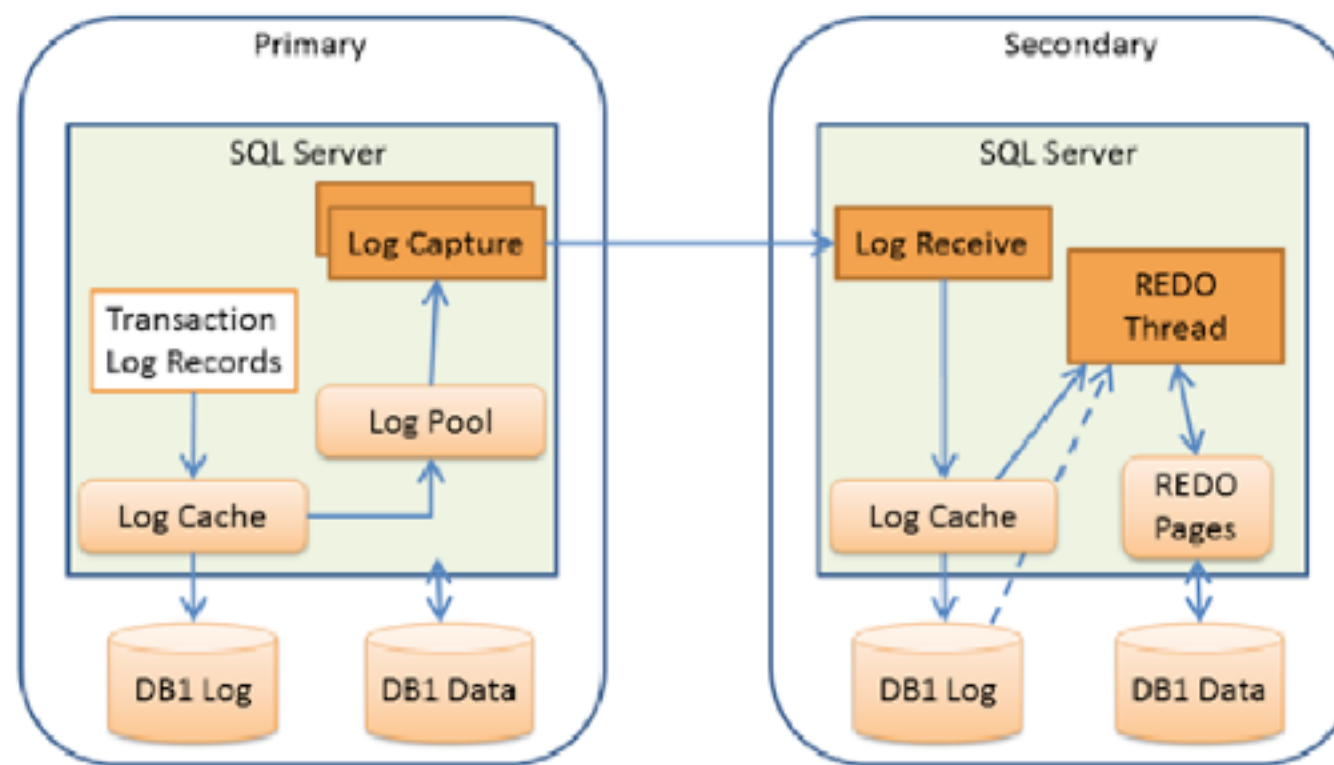
# Why is this important?

- Recovery Objectives
  - Recovery Point Objective - RPO
  - Recovery Time Objective - RTO
- Availability
  - How much data can we lose?
  - How fast will the system fail over?
- Monitoring and Trending
  - Establish a baseline for analysis - are we meeting those objectives?
  - Impact on resources
- Ownership
  - All of the components are monitored by the DBA

# Data Movement In Availability Groups

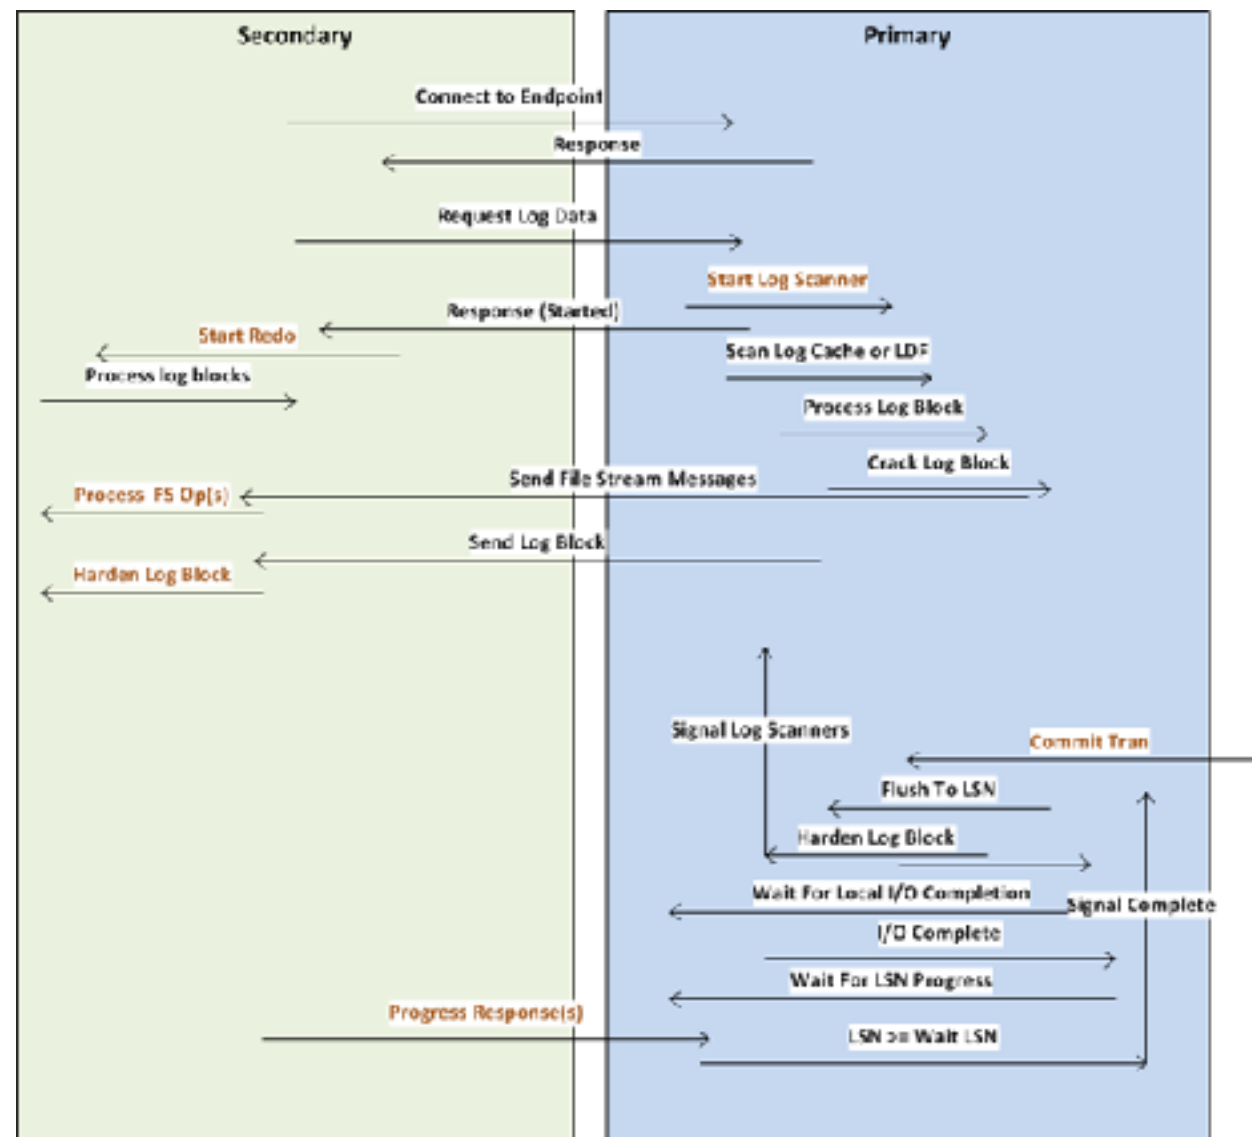
- Transaction log blocks are replicated to secondaries
- Replication mode
  - Synchronous
  - Asynchronous
- Database mirroring endpoint

# Data Movement in Availability Groups



From: SQLCAT's Guide to High Availability Disaster Recovery - <http://bit.ly/1U0vsss>

# Synchronous Commit



From: CSS SQL Engineers Blog - <http://bit.ly/1P8SzNc>

# Network Based Replication

- Strong working relationship with network team
  - Maintenance - patching, network outages, database
- Network conditions can impact your AG's availability
  - Latency - how **long** it takes for a packet of data to traverse the network from source to destination.
  - Bandwidth - how **much** data can be moved in a time interval



# Network Latency

- Often measured in milliseconds, sometimes microseconds
- Directly impacts network throughput
- TCP Congestion Window
- ping isn't your best measure of latency, by default it doesn't include any load...measure **your** workload
- It's often up to us to **PROVE** to the network team there is an issue
  - **Pinging 192.168.2.1 with 32 bytes of data:**
  - **Reply from 192.168.2.1: bytes=32 time=1001ms TTL=128**
- In synchronous mode, you have to wait for network latency

# Availability Group Flow Control

- Used in response to network and system conditions
- Log blocks exchange sequence numbers
- The AG will enter flow control mode IF:
  - The **primary** detects too many unacknowledged messages, the primary stops sending messages
  - The **secondary** needs to tell the primary to back off, likely due to resource constraints, it will send a flow control message to the primary to back off
- Primary polls every 1000ms for a change in flow control state
- Secondary will message primary to leave flow control mode

From: SQL Server PFE Blog - <http://bit.ly/1ZpGyIL>

# Database Synchronization States

- Not synchronizing
- Synchronized
- Synchronizing
- Reverting
- Initializing

<https://msdn.microsoft.com/en-us/library/ff877972.aspx>


# Failover Modes

- Automatic
  - Synchronous mode only
  - Synchronization state must be synchronized
  - Commonly used within a data center
- Manual
  - Synchronous or Asynchronous
  - Commonly used between data centers

<https://msdn.microsoft.com/en-us/library/hh213151.aspx>

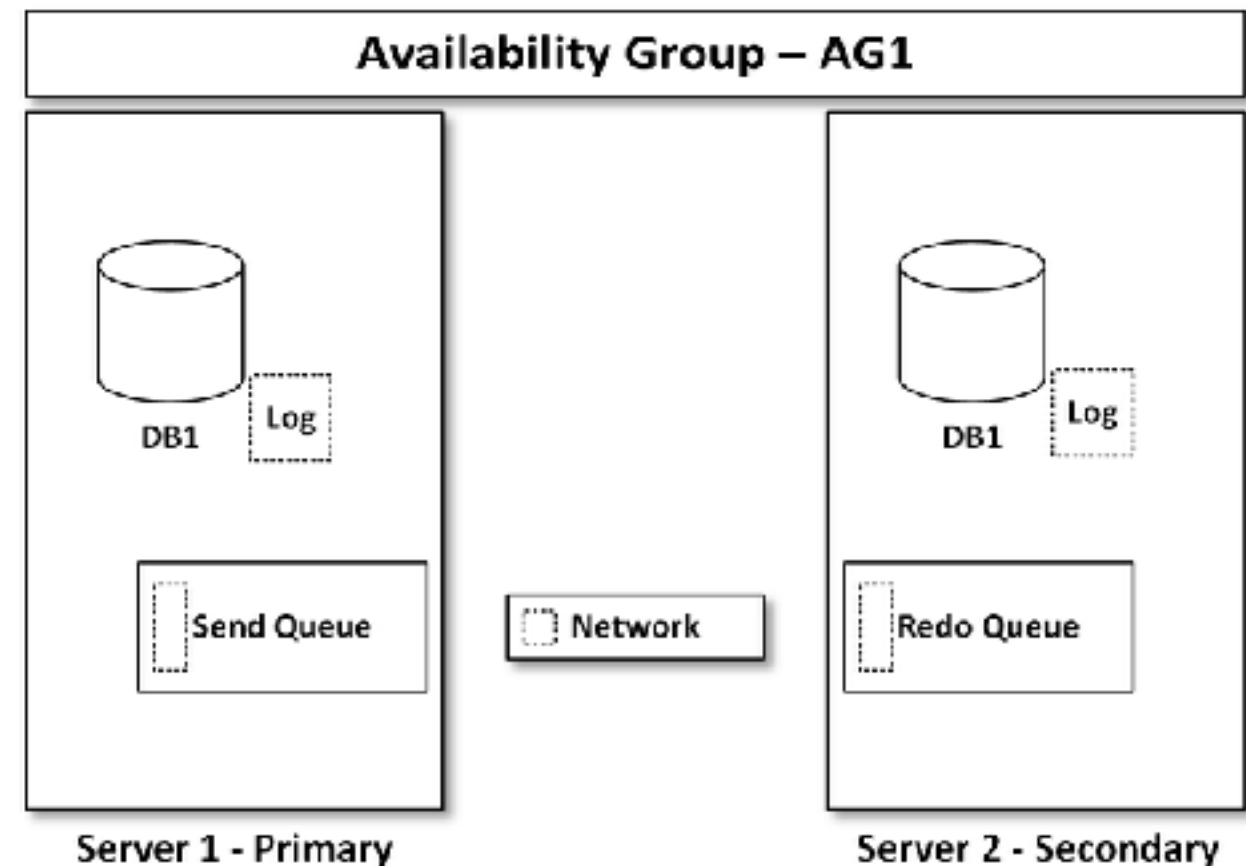
# Send Queue

- Queues log blocks to be sent to the secondaries
- Each replica maintains it's own view of the send queue
- Queued data is at risk to data loss in the event of a primary failure
- The send queue can grow due to an unreachable secondary, network outage, network latency and large amount of data change



# Redo Queue

- Queues log blocks received on the secondary
- Each replica has its own redo queue
- On failover, the redo queue must be completely processed
- The redo queue can grow due to a slow disk subsystem or resource contention or sustained outage and subsequent reconnection of a secondary



# Send Queue Impact on Availability


- When log generation on primary exceeds the rate they can be sent to the secondaries...
  - No automatic failover
  - Data loss
  - Stale data for reporting from secondaries
  - Stale data for off-loaded backups on secondaries
  - Off-loaded log backups can fail
  - Transaction delay
  - Fill up transaction logs
- **Even in synchronous mode!**

# Redo Queue Impact on Availability

- When log blocks received on the secondary exceed the rate they can be processed by the redo thread...
  - Delayed failover
    - Detect failure
    - Process Redo Queue
    - Crash recover database
  - Stale data for reporting from secondaries
  - Stale data for off-loaded backups on secondaries
  - Off-loaded log backups can fail
  - Transaction delay



# Log Backups



# Disconnected Replica

- When a synchronous secondary exceeds its session time-out, it changes to asynchronous commit mode
- When the secondary comes back online, it will attempt to re-sync, and resume synchronous mode
- WRITELOG - while replica is down
- HADR\_SYNCHRONIZING\_THROTTLE - time to go from SYNCHRONIZING to SYNCHRONIZED

# Transaction Delay

- In synchronous mode, when secondaries are behind, queries on the primary can be delayed
- HADR\_SYNC\_COMMIT
- HADR\_SYNCHRONIZING\_THROTTLE - replica back online

# Maintenance Events That Can Impact Availability

- Bulk data modifications
  - Database maintenance
  - Network or server maintenance
- 
- Carefully plan maintenance
  - Collaborate with other teams!

# Monitoring AG Performance

- Dynamic Management Views
  - `sys.dm_hadr_database_replica_states`
- Perfmon Counters
  - SQL Server:Availability Replica
    - Replication data - messages sent, bytes sent, flow control
  - SQL Server:Database Replica
    - Database data - log bytes sent, queue sizes, transaction delay per database

# Measuring Replication Latency

- `sys.dm_hadr_database_replica_states`
  - `log_send_queue_size`
  - `log_send_rate`
  - `redo_queue_size`
  - `redo_queue_rate`
- On the primary there's a row for each database on **each** replica
- On the secondaries there's a row for each database on **that** replica
- Replicas track their own values
- When a replica goes offline...
  - `log_send_queue_size` changes to NULL

log send queue is from primary to secondary

# Measuring Replication Latency - **ugh!!!**

- **Well, it looks like `sys.dm_hadr_database_replica_states` doesn't report the correct values for `log_send_rate` and `redo_queue_rate`**
  - Reported on Connect
    - <https://connect.microsoft.com/SQLServer/Feedback/Details/928582>
  - Known bug in SQL Server 2012 or 2014
    - <https://support.microsoft.com/en-us/kb/3012182>
    - Cumulative Update 5 or better
  - Observed in SQL 2016
  - Perfmon!

# Measuring Latency with Perfmon

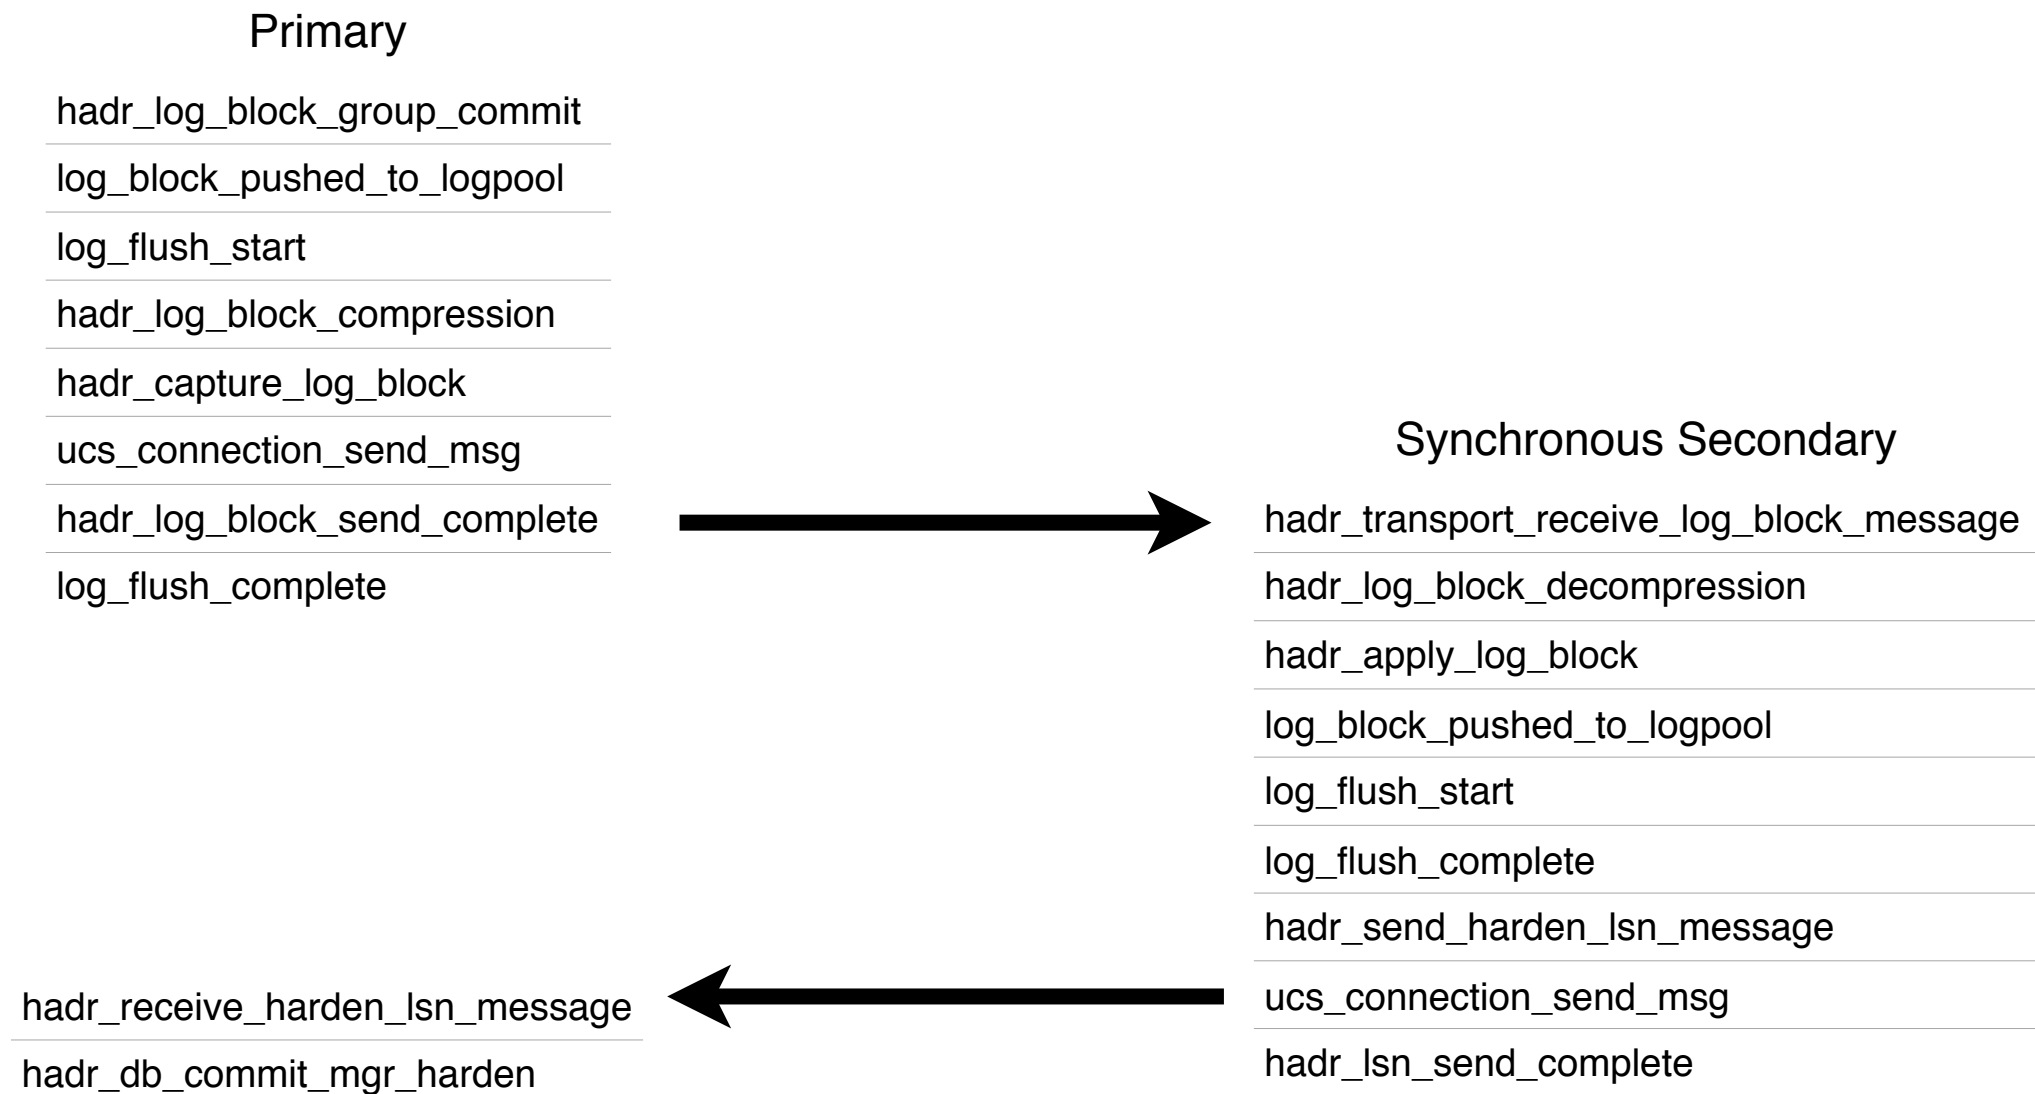
- Primary
  - SQLServer:Databases - Log Bytes Flushed/sec
  - SQLServer:Availability Replica - Bytes Sent to Replica/sec (compressed)
  - Network Interface - Bytes Sent/sec
- Secondaries
  - SQLServer:Availability Replica - Bytes Received From Replica/sec (compressed)
  - SQLServer:Database Replica - Log Bytes Received/sec (log send rate/uncompressed)
  - SQLServer:Database Replica - Redone Bytes/sec (log redo rate)
  - Network Interface - Bytes Received/sec



# Measuring Latency with Extended Events

**In SQL 2014 - SP2/2016 - SP1**

Each event has a measure duration



# Wait stats - sync vs. async

## Synchronous - HADR\_SYNC\_COMMIT

```
sqldk.dll!XeSosPkg::wait_info::Publish+0x138
sqldk.dll!SOS_Scheduler::UpdateWaitTimeStats+0x2bc
sqldk.dll!SOS_Task::PostWait+0x9e
sqlmin.dll!EventInternal<SuspendQueueSLock>::Wait+0x1fb
sqlmin.dll!SequencedObject<LogBlockId,SequencedWaitInfo<LogBlockId>0>::WaitUntilSequenceAdvances+0x160
sqlmin.dll!HaDrCommitMgr::HardenNotifyInternal+0x1af
sqlmin.dll!HaDrCommitMgr::HardenNotify+0xac
sqlmin.dll!RecoveryUnit::NotifyHardenParticipants+0x1e2
sqlmin.dll!RecoveryUnit::HardenLog+0x217
sqlmin.dll!XdesRMFull::CommitInternal+0x6b8
sqlmin.dll!XactRM::SinglePhaseCommit+0x1a1
sqlmin.dll!XactRM::CommitInternal+0x472
...
```


## Asynchronous - WRITELOG

```
sqldk.dll!XeSosPkg::wait_info::Publish+0x138
sqldk.dll!SOS_Scheduler::UpdateWaitTimeStats+0x2bc
sqldk.dll!SOS_Task::PostWait+0x9e
sqlmin.dll!SQLServerLogMgr::WaitLCFlush+0x219
sqlmin.dll!SQLServerLogMgr::LogFlush+0x29e
sqlmin.dll!SQLServerLogMgr::WaitLogWritten+0x17
sqlmin.dll!RecoveryUnit::HardenLog+0x25e
sqlmin.dll!XdesRMFull::CommitInternal+0x6b8
sqlmin.dll!XactRM::SinglePhaseCommit+0x1a1
sqlmin.dll!XactRM::CommitInternal+0x472
...
```

# Monitoring Tools


- Build your own
- AlwaysOn Dashboard
- Third Party Tool
  - SentryOne SQL Sentry for SQL Server
  - Redgate SQL Monitor

# Demo




# Demo

# Demo



# Real World Example



# Dealing With Slow Replication Latency

- Identify your bottleneck and mitigate it
  - Minimize log generation
    - Use smart index maintenance/Better Indexes
  - More bandwidth
    - Perhaps a dedicated network connection
  - Better hardware
    - Log throughput on secondaries needs to be equal to primary
- Upgrade SQL Server
  - 2012/2014 single threaded redo - ~45MB/sec
  - 2016 multi-threaded redo - ~600MB/sec



# Key Takeaways

- It is imperative to track and trend replication latency in your Availability Groups so you can answer the questions
  - How much data can I lose?
  - How long it will take to failover?
- Monitor and trend `send_queue` and `redo_queue` in `sys.dm_hadr_database_replica_states` on replicas to measure availability impact
- Understand how much log is generated in your databases
- Understand your system's operations, consider downtime for patching and network maintenance

# Key Takeaways

- Plan database maintenance
- Use a smart index maintenance strategy!
- Offloaded backups
  - If availability is most important, backup on primary

# Need more data or help?

**<http://www.centinosystems.com/blog/talks/>**

Links to resources

Demos

Presentation

**[aen@centinosystems.com](mailto:aen@centinosystems.com)**

**[@nocentino](#)**

**[www.centinosystems.com](http://www.centinosystems.com)**

**Solving tough business challenges with technical innovation**



# Questions?

# References

- <http://www.centinosystems.com/blog/sql/designing-for-offloaded-backups-in-alwayson-availability-groups/>
- <http://www.centinosystems.com/blog/sql/designing-for-offloaded-log-backups-in-alwayson-availability-groups-monitoring/>
- <http://www.centinosystems.com/blog/sql/monitoring-availability-groups-with-redgates-sql-monitor>
- <https://msdn.microsoft.com/en-us/library/ff878537.aspx>
- <https://msdn.microsoft.com/en-us/library/ff877972.aspx>