

Deploying SQL Server in Kubernetes

Anthony E. Nocentino
aen@centinosystems.com



Anthony E. Nocentino

- **Consultant and Trainer**
- **Founder and President of Centino Systems**
 - Specialize in system architecture and performance
 - Masters Computer Science
 - Microsoft MVP - Data Platform - 2017 - 2020
 - Linux Foundation Certified Engineer
 - Friend of Redgate - 2015-2020
- **email:** aen@centinosystems.com
- **Twitter:** @nocentino
- **Blog:** www.centinosystems.com/blog
- **Pluralsight Author:** www.pluralsight.com



Agenda

- **Deploying SQL Server in Kubernetes**
 - Data Persistency and Storage in Kubernetes
 - Pod Configuration and Running SQL Server in a Pod
 - Disk and Resource Configurations
 - Backups
 - The ~~Future~~ **Present** of SQL Server and Kubernetes

Kubernetes 101


- Container Orchestrator
- Pods are Container Based Applications
- A Cluster is a Collection of Compute Resources
- Infrastructure Abstraction
- Declarative Configuration in Code
- Desired State



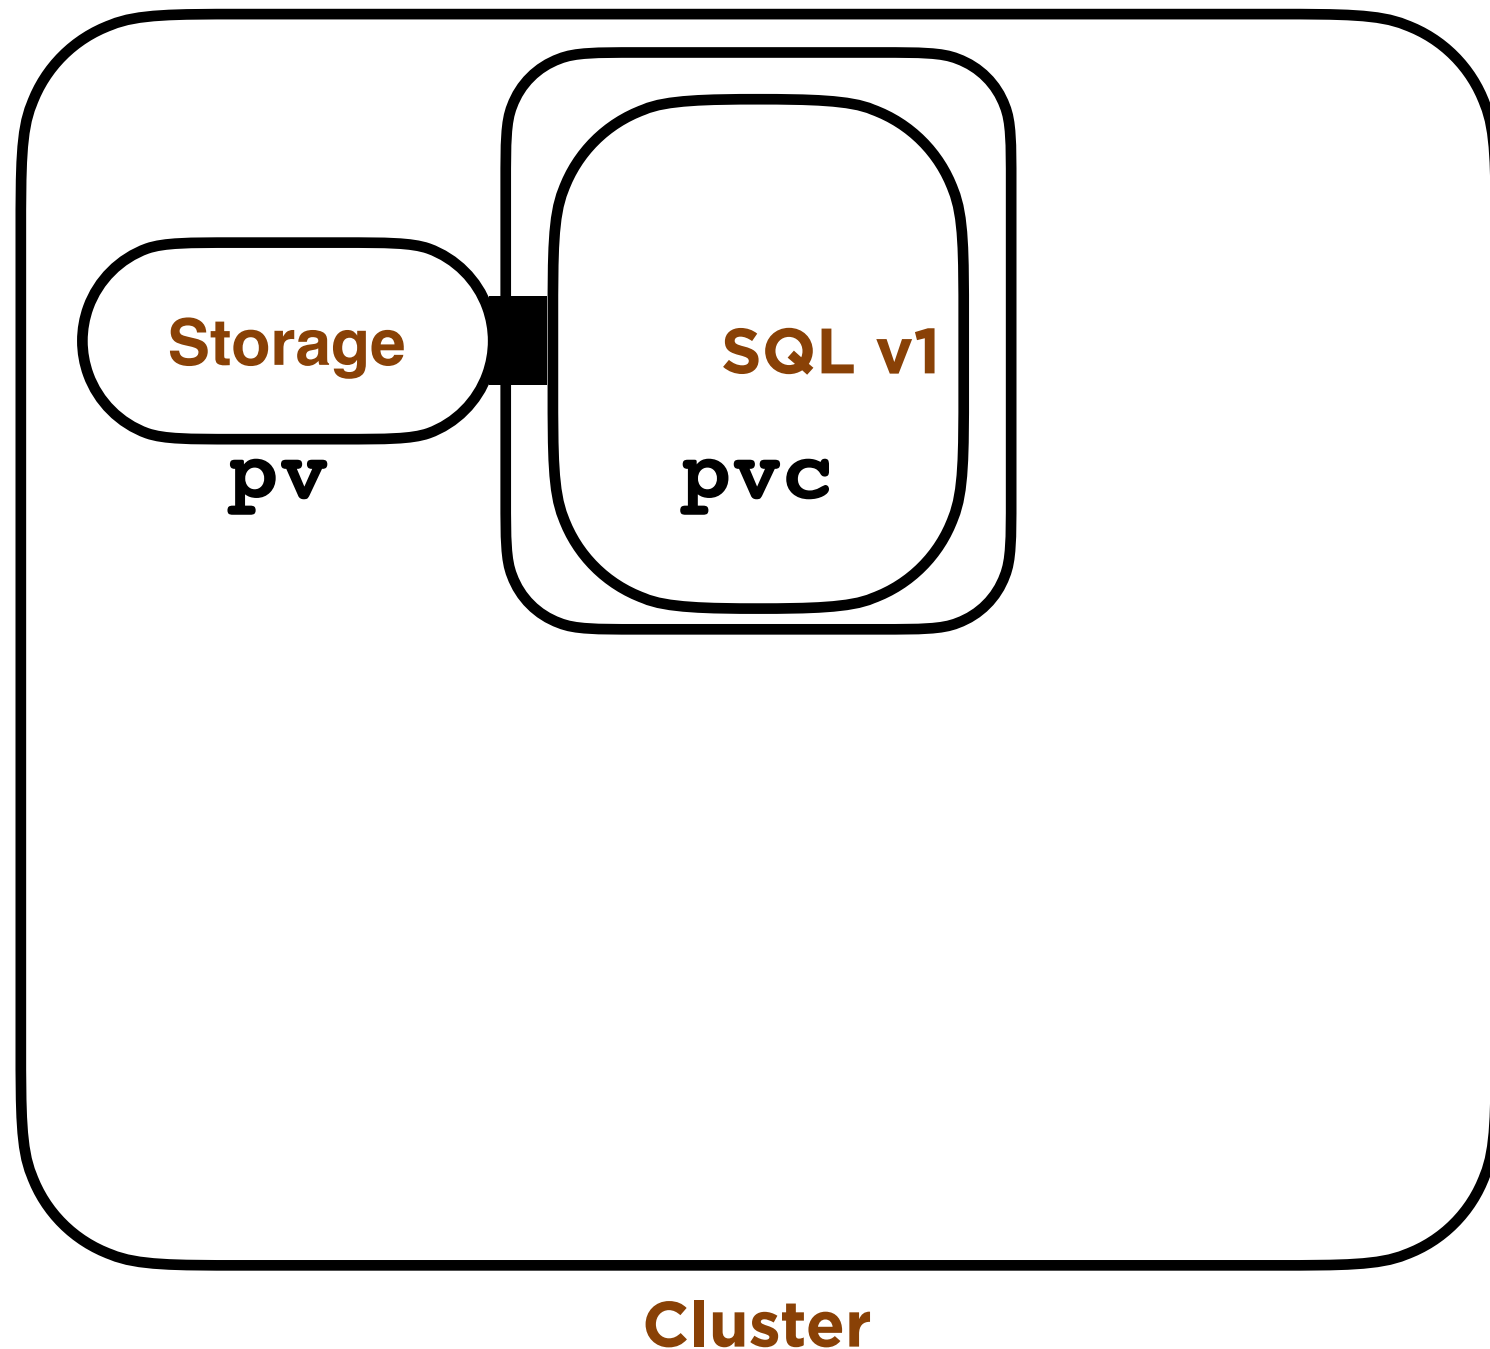
Running SQL Server in Kubernetes

- A Pod goes back its initial state each time it's deployed
- **State** - where do we store data?
- **Configuration** - how do we configure SQL Server?

Decoupling Data and Computation




Storage in Kubernetes



- **Persistent Volumes (pv)**
 - Administrator defined storage
 - iSCSI, NFS, FC, AzureDisk...many more
- **Persistent Volume Claims (pvc)**
 - The Pod “claims” the **pvc**
 - The **pvc** is mapped to the **pv** by k8s
 - Decouples the Pod and the storage

Data Persistency in SQL Server in K8S

- Define **Persistent Volumes/Persistent Volume Claims**
 - Instance directory (error log, default trace, etc..)
 - **`/var/opt/mssql/`**
 - User Database default directory
 - **`/var/opt/mssql/data`**



Defining Persistent Volumes and Persistent Volume Claims

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv-nfs-data
  labels:
    disk: data
spec:
  capacity:
    storage: 10Gi
  accessModes:
    - ReadWriteOnce
  persistentVolumeReclaimPolicy: Retain
  nfs:
    server: 172.16.94.5
    path: "/export/volumes/sql/data"
```

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pvc-nfs-data
spec:
  selector:
    matchLabels:
      disk: data
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 10Gi
```

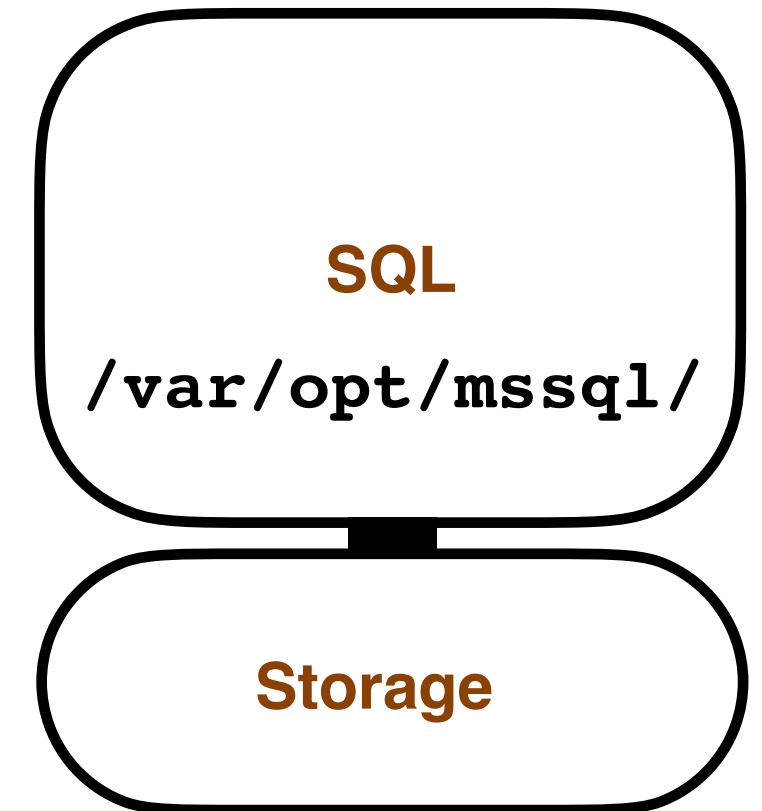
Configuring SQL Server in a Pod

- In our Pod configuration we define **Environment Variables**
 - Used at initial startup to configure the SQL Instance
 - **ACCEPT_EULA**
 - **MSSQL_SA_PASSWORD**
 - Stored in the cluster as a **Secret**
 - Pods go back their initial state of the container image on creation
 - But some settings are persisted in master, right...yep!

<https://docs.microsoft.com/en-us/sql/linux/sql-server-linux-configure-environment-variables>

Running SQL Server in a Pod (con't)

- In our Pod configuration define our storage configuration (**pvc**)
- Initial Pod deployment
 - If there's no system databases in the default data directory...
 - **/var/opt/mssql/data**
 - They're copied into the default data directory from the SFPs
- On subsequent Pod deployments the storage is attached into the 'new' Pod
 - Databases are already there
 - Master is read...contains our instance's configuration and state
 - Defined and accessible user databases are brought online



Define SQL Server in a Pod in YAML

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: mssql-deployment
```

```
spec:
```

```
  replicas: 1
```

```
  strategy:
```

```
    type: Recreate
```

```
  selector:
```

```
    matchLabels:
```

```
      app: mssql
```

```
spec:
```

```
  securityContext:
```

```
    fsGroup: 10001
```

```
  containers:
```

```
  - name: mssql
```

```
    image: '.../mssql/server:2019-CU5-ubuntu-18.04'
```

```
    ports:
```

```
    - containerPort: 1433
```

```
    env:
```

```
    - name: ACCEPT_EULA
```

```
      value: "Y"
```

```
    - name: SA_PASSWORD
```

```
      valueFrom:
```

```
        secretKeyRef:
```

```
          name: mssql
```

```
          key: SA_PASSWORD
```

```
  volumeMounts:
```

```
  - name: mssqldb
```

```
    mountPath: /var/opt/mssql
```

```
volumes:
```

```
  - name: mssqldb
```

```
    persistentVolumeClaim:
```

```
      claimName: pvc-sql-data
```



Advanced Disk Topologies for SQL Server

- Define your **Persistent Volumes** and **Persistent Volume Claims**
- Use environment variables to specify default directories on Pod at startup
 - **MSSQL_DATA_DIR (/data)**
 - **MSSQL_LOG_DIR (/log)**
- New user databases will be created in these locations
- On Pod creation
 - All **PV/PVCs** will be mounted in the container at the defined locations
 - Master will online the databases

Resource Management

- Pod level resource management

- CPU and Memory

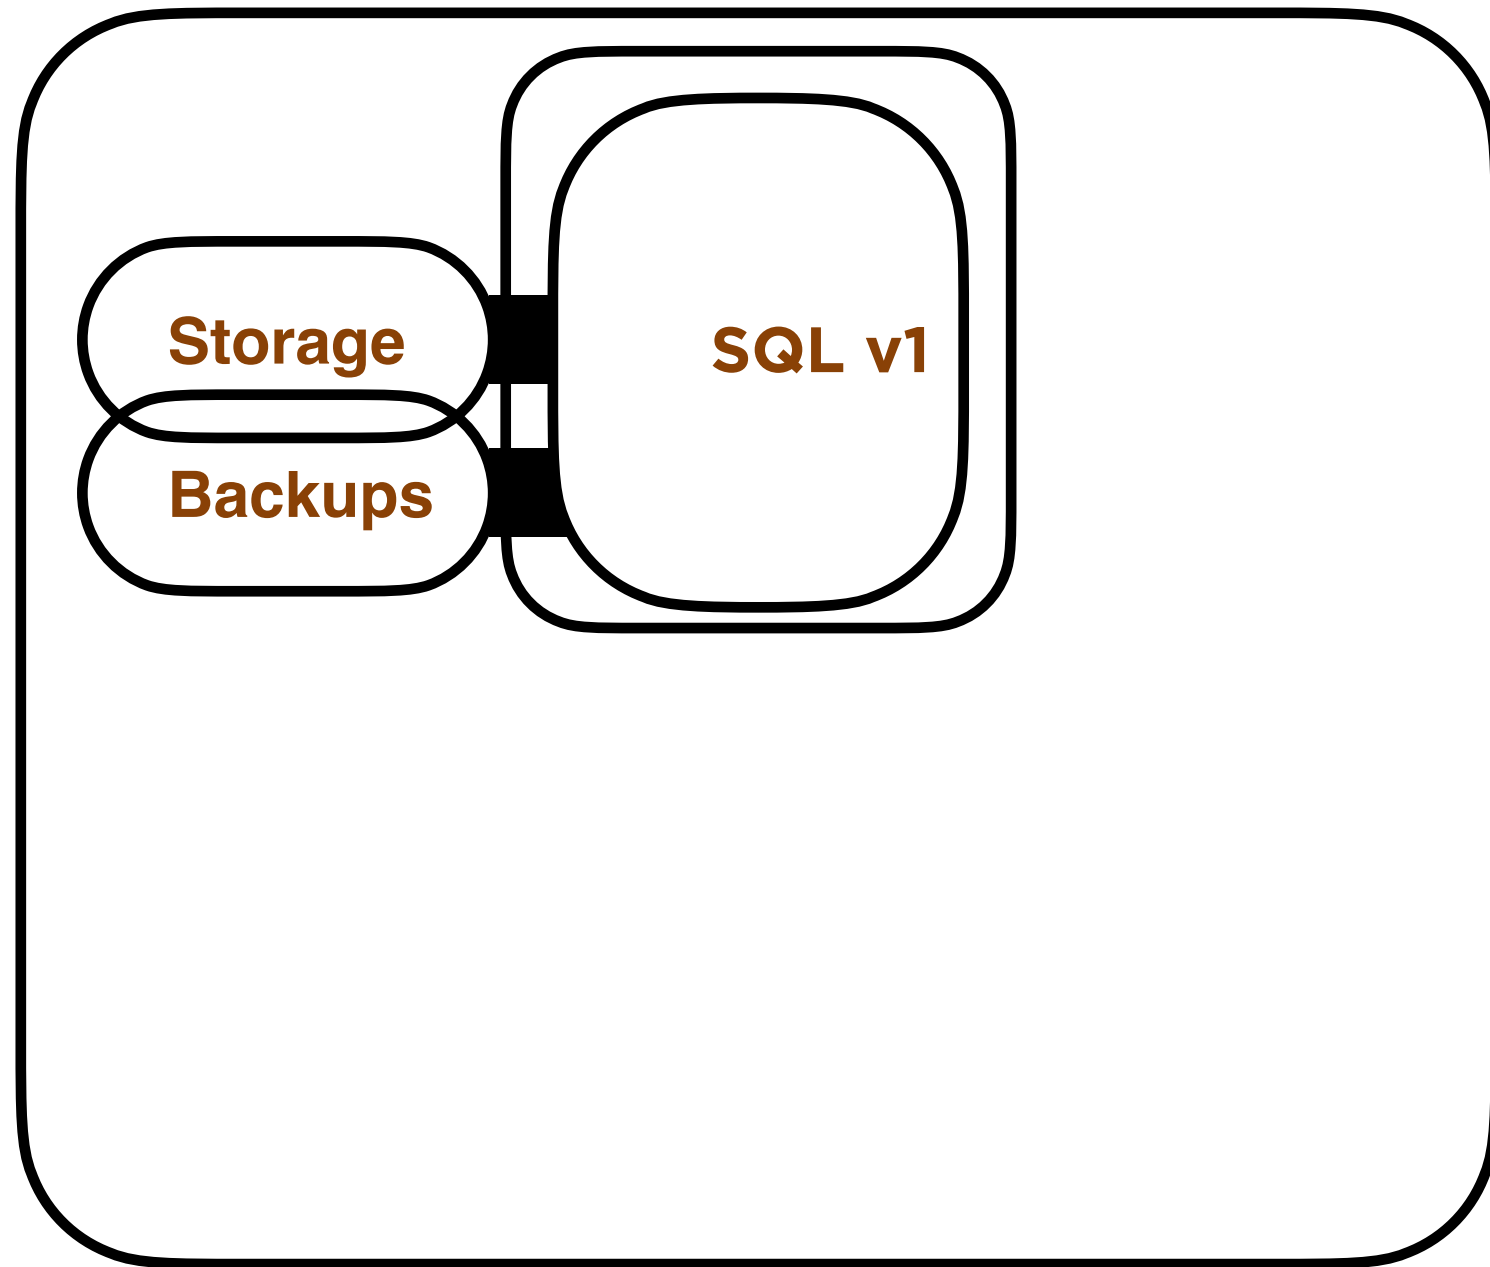
- **requests** - guaranteed
 - **limits** - upper limit
 - No limits by default
- ```
containers:
- name: mssql
 image: '.../mssql/server:2019-CU5-ubuntu-18.04'
 resources
 requests:
 cpu: 1
 memory: 1Gi
```

- Server Instance settings still apply
- Kind of like multi-instance clusters
- Workload is stopped and started when moved between Nodes


URL

AzureFiles

# Backups!



Cluster

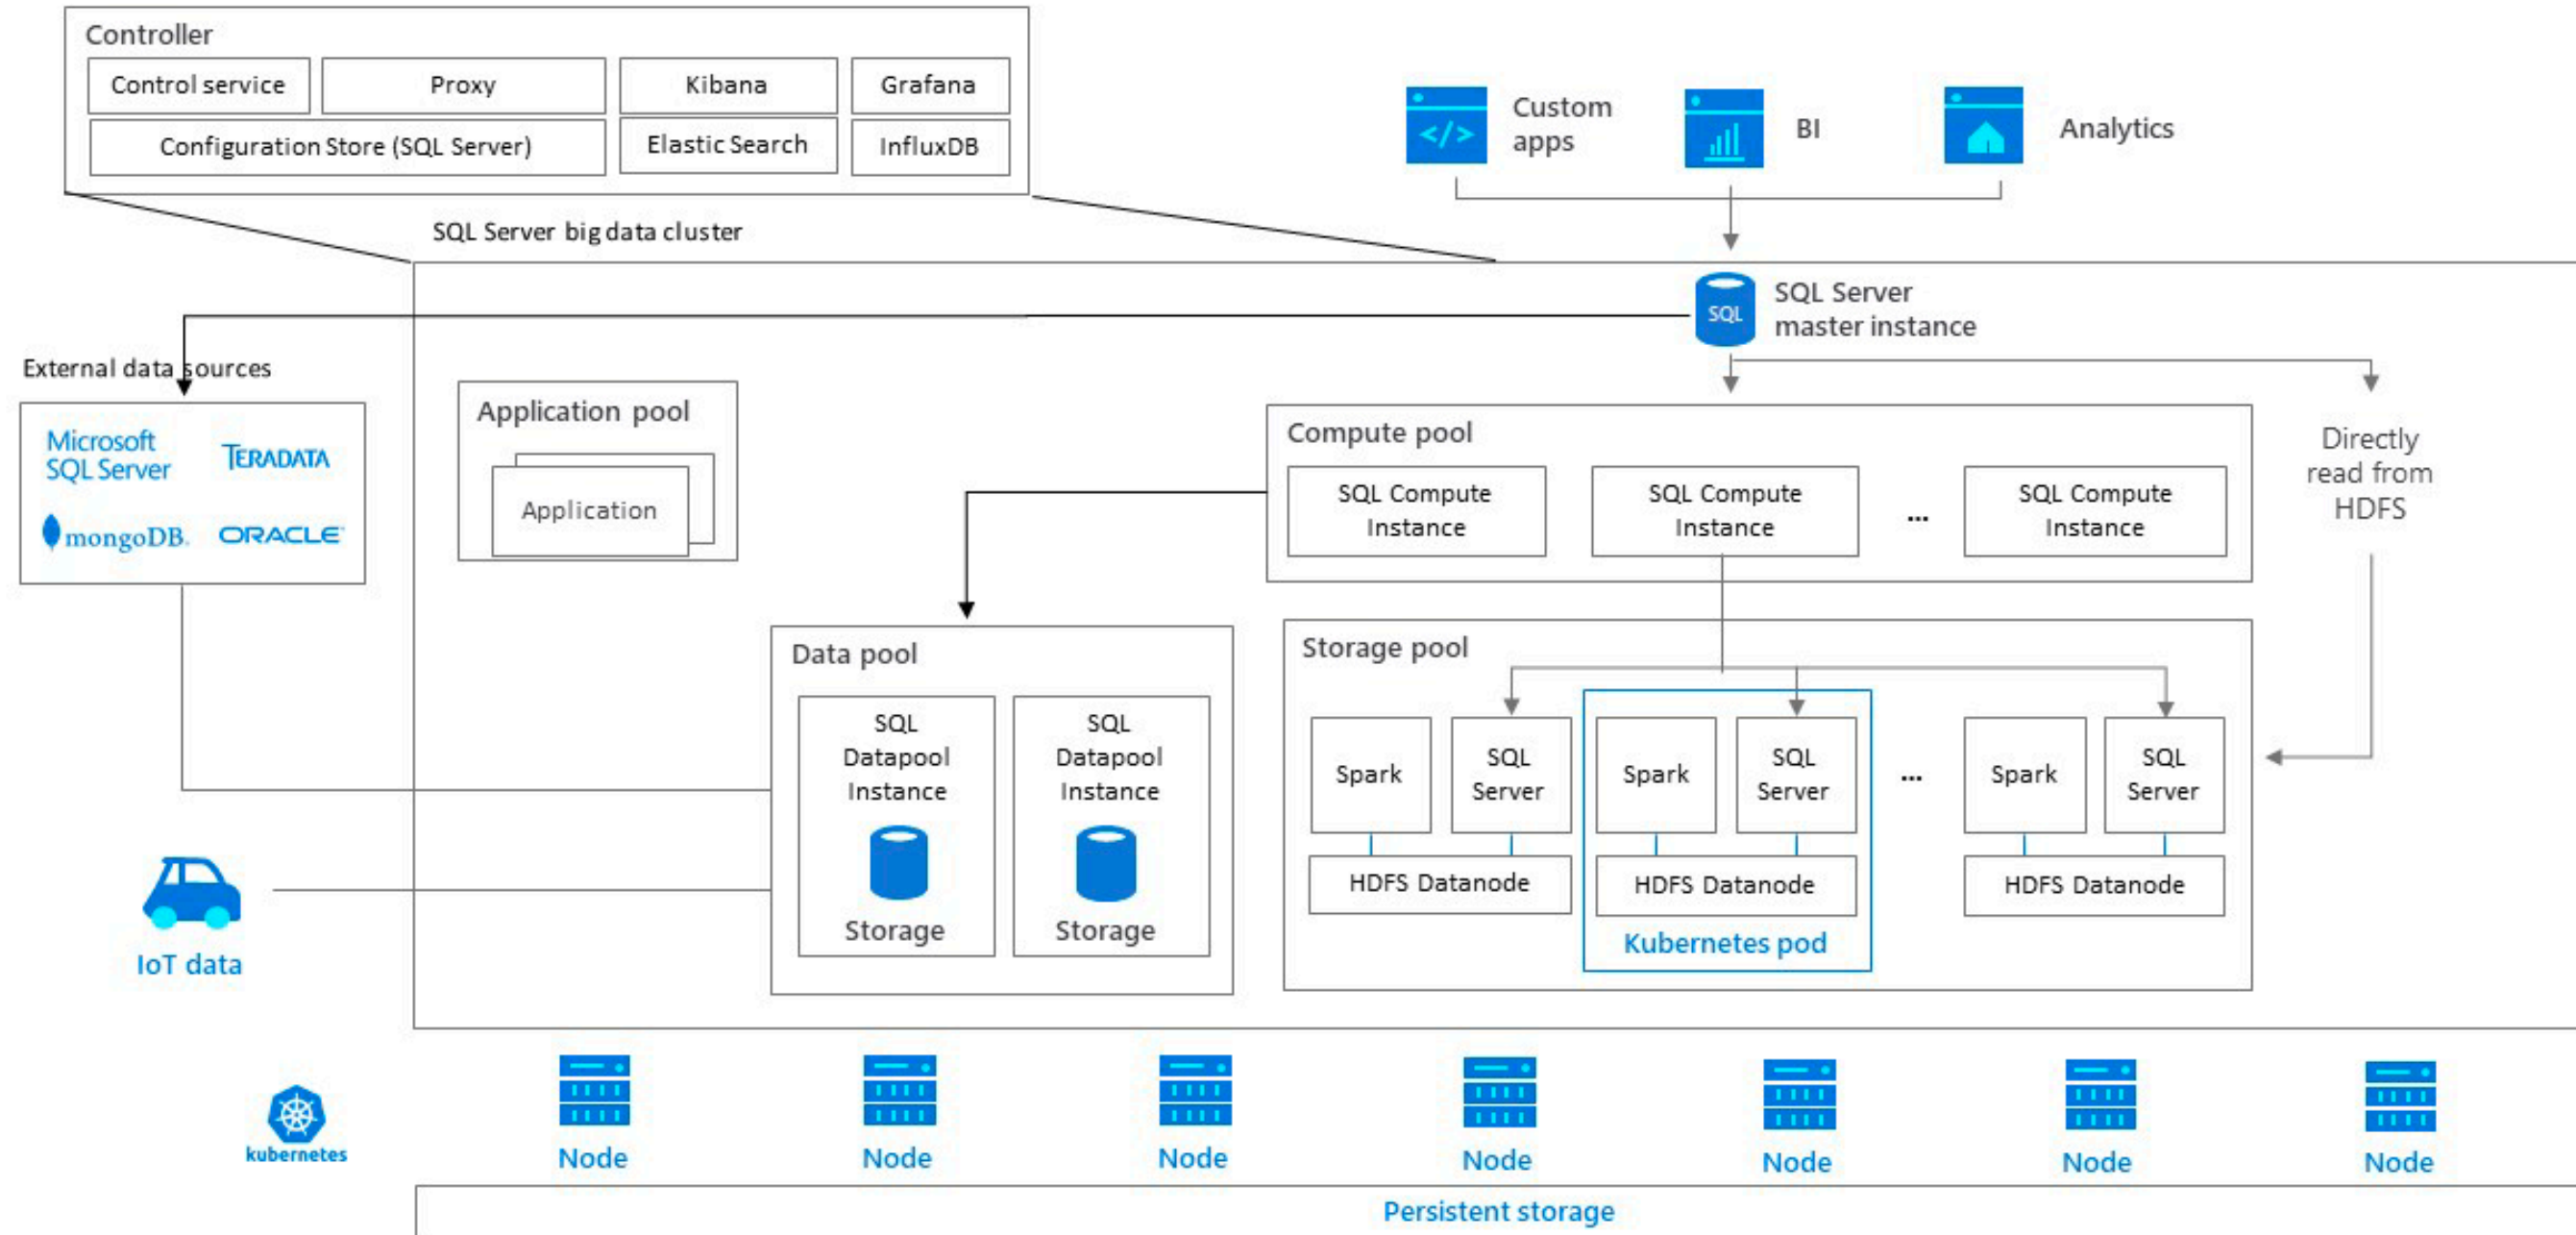
- Persistent Volume (Shared or Dedicated)
  - AzureDisk
  - AzureFile
  - NFS/iSCSI/FC
- To URL
- Drive the backup jobs with normal techniques
  - Ola Hallengren's
  - Maintenance Plans
  -  dbatools

# Demo!

- Deploying SQL Server in a **Deployment** with Persistent Storage
  - Disk Topology
  - Setting Resource Limits
  - High Availability
  - Backing up SQL Server in Kubernetes



# Big Data Clusters



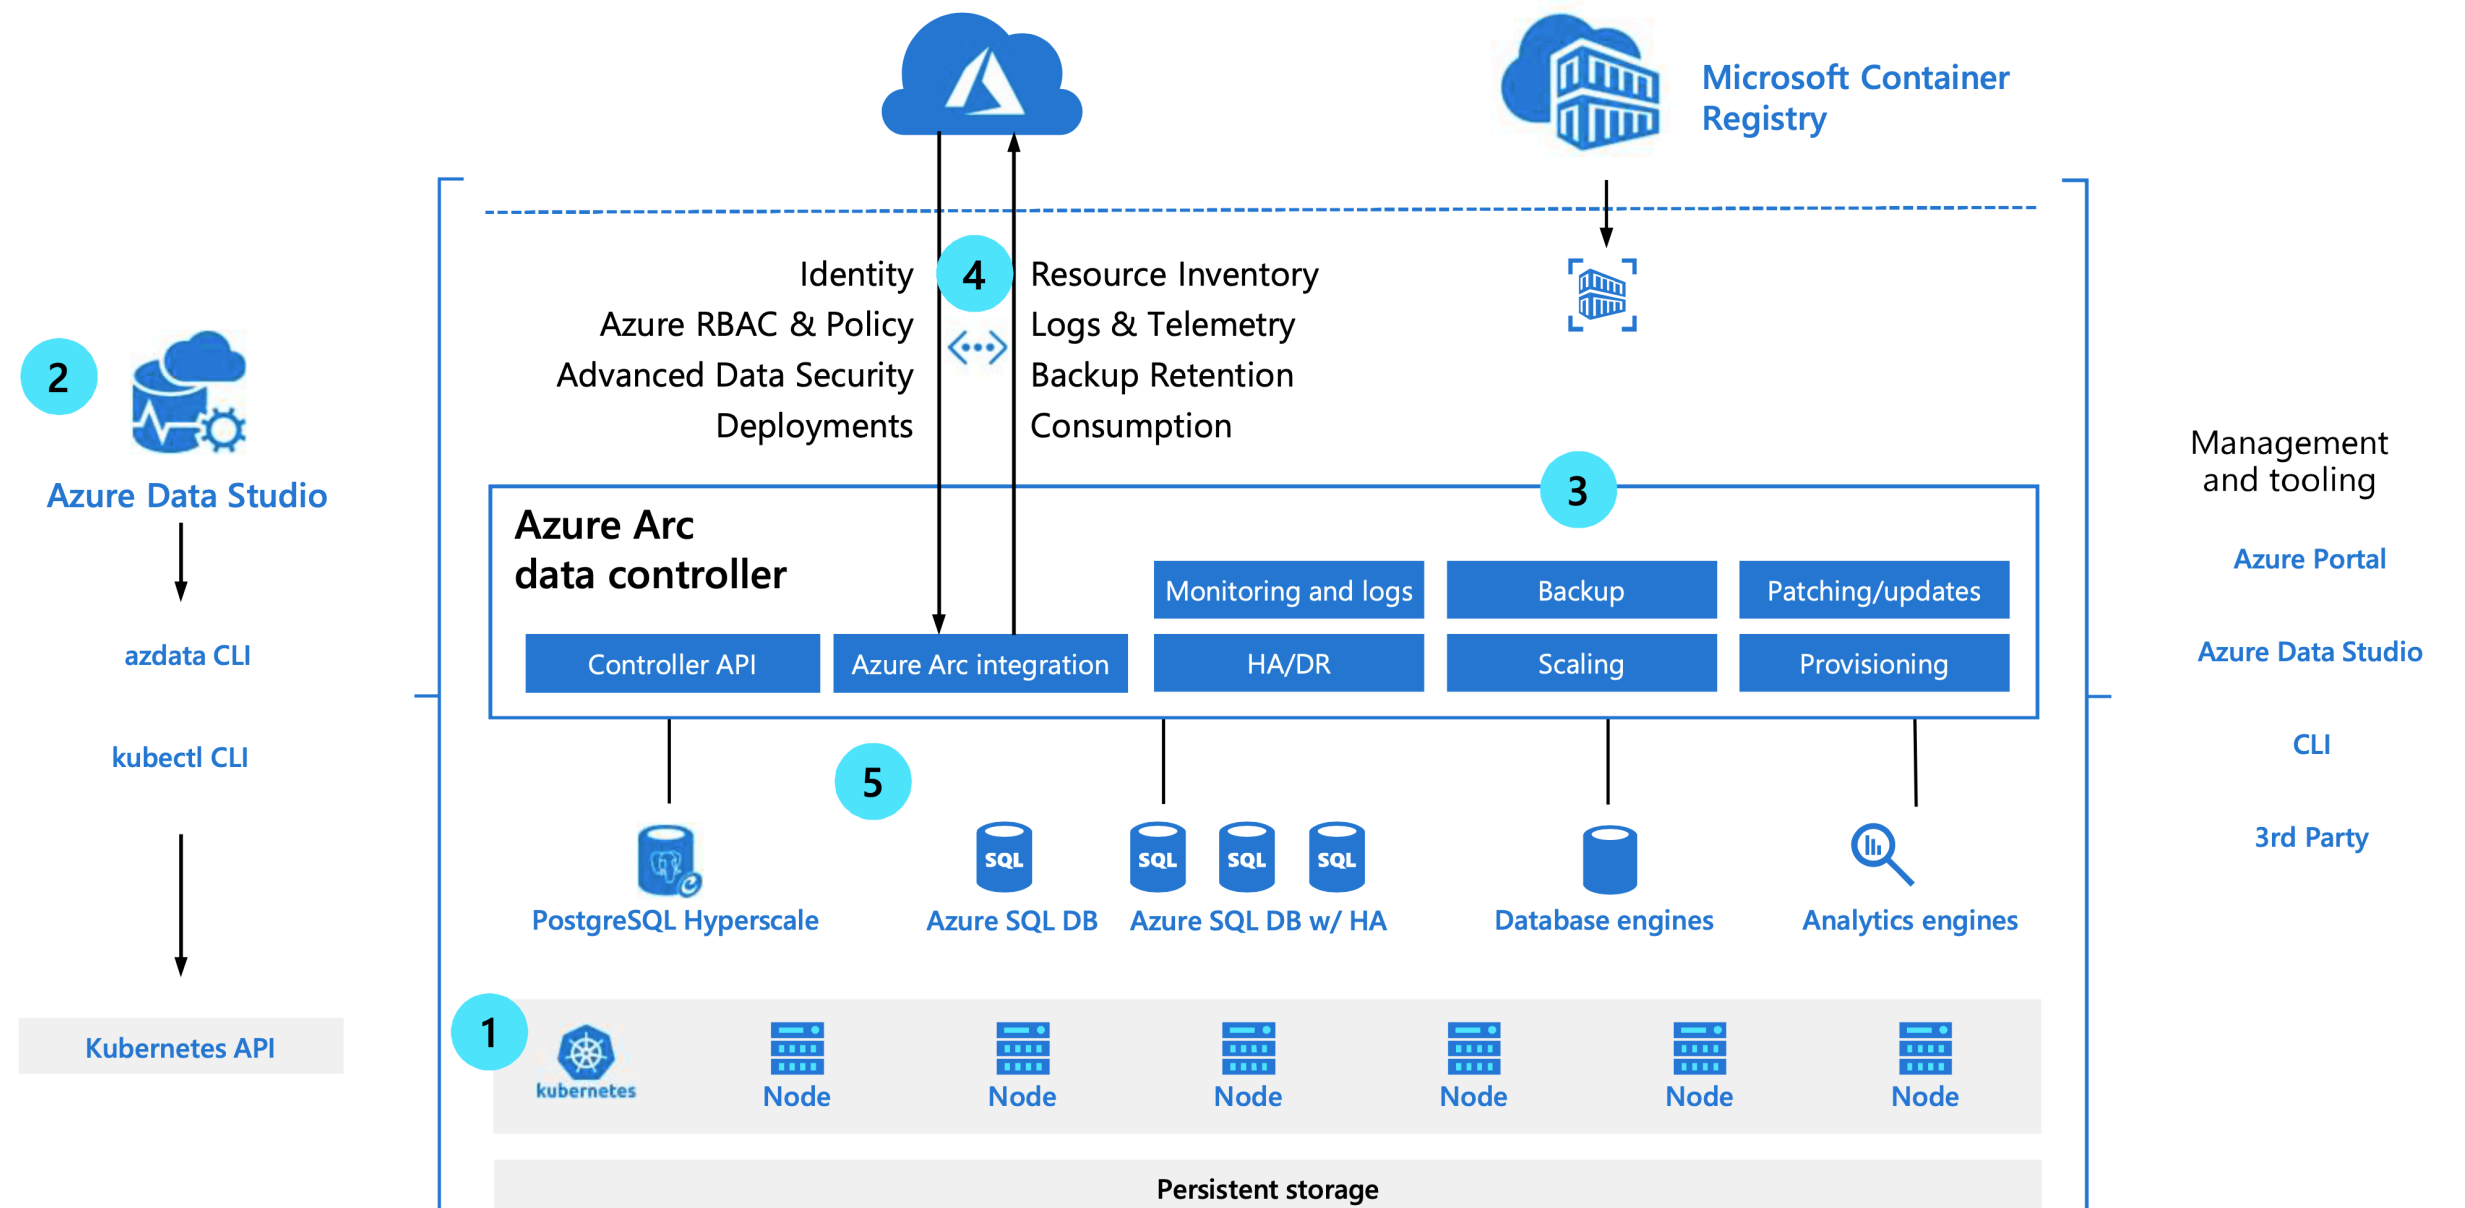
<https://docs.microsoft.com/en-us/sql/big-data-cluster/big-data-cluster-overview?view=sqlallproducts-allversions>

# Azure Arc - Data Services

## How it works: architecture of Azure data services on customer infrastructure

A few steps to get Azure data services in your environment:

- 1 Have Kubernetes on your infrastructure
- 2 Prepare environment with APIs and CLIs
- 3 Install Azure Arc data controller
- 4 Connect to Azure
- 5 Deploy and run Azure data services for your workloads



# Review

- **Deploying SQL Server in Kubernetes**
  - Data Persistency and Storage in Kubernetes
  - Pod Configuration and Running SQL Server in a Pod
  - Disk and Resource Configurations
  - Backups
  - The Future of SQL Server and Kubernetes

# More Resources

- **Docker for Windows/Mac**
- **Managed Service Providers**
  - Azure Kubernetes Service (**AKS**)
    - <https://docs.microsoft.com/en-us/azure/aks/kubernetes-walkthrough>
  - **Elastic Container Service for Kubernetes (EKS)**
    - <https://aws.amazon.com/getting-started/projects/deploy-kubernetes-app-amazon-eks/>
  - **Google Kubernetes Engine (GKE)**
    - <https://cloud.google.com/kubernetes-engine/docs/how-to/>
- **Pluralsight**
  - <https://app.pluralsight.com/profile/author/anthony-nocentino>